Setting Up Apache Airflow with Docker

1. Install Docker Desktop

- Download and install Docker Desktop from: https://www.docker.com/products/docker-desktop
- After installation, verify Docker is working: docker --version docker-compose --version

You should see version numbers for both.

```
C:\Users\nithi>docker --version
Docker version 28.3.2, build 578ccf6
```

C:\Users\nithi>docker-compose --version
Docker Compose version v2.39.1-desktop.1

2. Create Airflow Project Directory

In PowerShell or VS Code terminal: mkdir airflow cd airflow

This will be the working directory for Airflow.

3. Create . env File

Inside C:\airflow create a file named .env and add: AIRFLOW_IMAGE_NAME=apache/airflow:2.10.2-python3.8 AIRFLOW_UID=50000

- AIRFLOW_IMAGE_NAME: Specifies Airflow image version.
- AIRFLOW UID: Ensures correct file permissions on Windows.

4. Create docker-compose.yaml

Inside C:\airflow, create docker-compose.yaml with the following content:

```
# Licensed to the Apache Software Foundation (ASF) under one
# or more contributor license agreements. See the NOTICE file
# distributed with this work for additional information
# regarding copyright ownership. The ASF licenses this file
# to you under the Apache License, Version 2.0 (the
# "License"); you may not use this file except in compliance
# with the License. You may obtain a copy of the License at
#
# http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing,
# software distributed under the License is distributed on an
# "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY
# KIND, either express or implied. See the License for the
# specific language governing permissions and limitations
# under the License.
#
```

```
x-airflow-common:
 &airflow-common
you can use your extended image.
you placed the docker-compose.yaml
   &airflow-common-env
airflow.providers.fab.auth manager.fab auth manager.FabAuthManager
postgresql+psycopg2://airflow:airflow@postgres/airflow
db+postgresql://airflow:airflow@postgres/airflow
```

```
usage) build/extend Airflow image.
in the local config folder
   AIRFLOW CONFIG: '/opt/airflow/config/airflow.cfg'
 user: "${AIRFLOW UID:-50000}:0"
   &airflow-common-depends-on
     POSTGRES USER: airflow
     POSTGRES PASSWORD: airflow
```

```
redis:
    # Redis is limited to 7.2-bookworm due to licencing change
    # https://redis.io/blog/redis-adopts-dual-source-available-licensing/
image: redis:7.2-bookworm
    expose:
        - 6379
    healthcheck:
        test: ["CMD", "redis-cli", "ping"]
        interval: 10s
        timeout: 30s
        retries: 50
```

```
<<: *airflow-common
restart: always
 <<: *airflow-common-depends-on
<<: *airflow-common
 <<: *airflow-common-depends-on
<<: *airflow-common
 <<: *airflow-common-depends-on
```

```
<<: *airflow-common
       - "CMD-SHELL"
'celery@$${HOSTNAME}" || celery --app
     <<: *airflow-common-env
    <<: *airflow-common-depends-on
    airflow-webserver:
   <<: *airflow-common
    <<: *airflow-common-depends-on
   <<: *airflow-common
```

```
echo -e "\033[1;33mWARNING!!!: AIRFLOW UID not set!\e[0m"
below to set "
owned by root."
with manually created .env file:"
       one meg=1048576
          echo -e "\033[1;33mWARNING!!!: Not enough CPUS available for
```

```
$$((disk available * 1024 )))"
run Airflow (see above)!\e[0m"
resources available:"
       echo "Airflow version:"
        echo "Files in shared volumes:"
missing."
       echo "Files in shared volumes:"
        echo "Change ownership of files in /opt/airflow to
${AIRFLOW UID}:0"
${AIRFLOW UID}:0"
```

```
echo
echo "Files in shared volumes:"
echo
ls -la /opt/airflow/{logs,dags,plugins,config}
```

```
<<: *airflow-common-env
   <<: *airflow-common
     <<: *airflow-common-env
     - airflow
     <<: *airflow-common-depends-on
docker-compose --profile flower up
   <<: *airflow-common
     - flower
```

```
restart: always
  depends_on:
          <<: *airflow-common-depends-on
          airflow-init:
          condition: service_completed_successfully
volumes:
    postgres-db-volume:</pre>
```

5. Initialize Airflow

Before starting Airflow, initialize the database:

docker-compose up airflow-init

```
C:\airflow>docker-compose up airflow-init

[+] Running 22/22

/ airflow-init Pulled

[+] Running 3/3

/ Container airflow-redis-1

Attaching to airflow-airflow-init-1

airflow-init-1

| WARNING!!!: Not enough memory available for Docker.
```

6. Start Airflow

Run containers in the background:

docker-compose up -d

```
C:\airflowJocker-compose up -d

[+] Running 8/8r Desktop View Config

Vontainer airflow-redis-1

Container airflow-postgres-1

Container airflow-airflow-triggerer-1

Container airflow-airflow-dag-processor-1

Container airflow-airflow-webserver-1

Container airflow-airflow-scheduler-1

Container airflow-airflow-scheduler-1

Container airflow-airflow-scheduler-1

Container airflow-airflow-scheduler-1

Container airflow-airflow-oscheduler-1

Started

Started
```

Check status:

docker ps

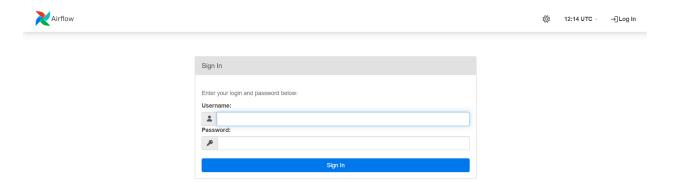
```
C:\airflow>docker ps
CONTAINER ID IMAGE
                                                         COMMAND
                                                                                       CREATED
                                                                                                              STATUS
                                                                                                                                                   PORTS
           NAMES
                apache/airflow:2.10.2-python3.8
                                                         "/usr/bin/dumb-init ..."
                                                                                                              Up 35 minutes (healthy)
                                                                                      37 minutes ago
airflow-airflow-worker-1
1cb1ce426938 apache/airflow-2.10.2-python3.8
080/tcp airflow-airflow-webserver-1
                                                         "/usr/bin/dumb-init ..." 37 minutes ago
                                                                                                                                                   0.0.0.0:8080->8080/tcp, [::]:8080->8
                                                                                                              Up 36 minutes (healthy)
3a82c2717c00 apache/airflow:2.10.2-python3.8 "/usr/bin/dumb-init ..." 37 minutes ago airflow-airflow-scheduler-1
                                                                                                              Up 36 minutes (healthy)
                                                                                                                                                   8080/tcp
5188fcd5f630 apache/airflow:2.10.2-python3.8
airflow-airflow-dag-processor-1
f058304ab9ab apache/airflow:2.10.2-python3.8
                                                         "/usr/bin/dumb-init ..." 37 minutes ago
                                                                                                              Restarting (1) 4 seconds ago
                                                         "/usr/bin/dumb-init ..." 37 minutes ago
                                                                                                              Up 36 minutes (healthy)
                                                                                                                                                   8080/tcp
           airflow-airflow-triggerer-1
c8bd1c2e615b postgres:13
airflow-postgres-1
                                                          "docker-entrypoint.s..." About an hour ago Up About an hour (healthy)
                                                                                                                                                   5432/tcp
d957d9682184 redis:7.2-bookworm
                                                          "docker-entrypoint.s..." About an hour ago Up About an hour (healthy)
                                                                                                                                                   6379/tcp
           airflow-redis-1
C:\airflow>
```

You should see containers for webserver, scheduler, worker, etc.

7. Access Airflow UI

Open browser and go to:

http://localhost:8080



8. Create Admin User

Run the following command to create an **Admin user**:

docker-compose run airflow-worker airflow users create \

- --role Admin \
- --username admin \
- --email admin \
- --firstname admin \
- --lastname admin \
- --password admin

Now login at http://localhost:8080 with:

Username: adminPassword: admin

